

## 1. Co-rewarding: Stable Self-supervised RL for Eliciting Reasoning in Large Language Models

**Accession number:** 20250359596

**Authors:** Zhang, Zizhuo (1); Zhu, Jianing (1); Ge, Xinmu (2, 3); Zhao, Zihua (2); Zhou, Zhanke (1); Li, Xuan (1); Feng, Xiao (1); Yao, Jiangchao (2); Han, Bo (1)

**Author affiliation:** (1) TMLR Group, Department of Computer Science, Hong Kong Baptist University, Hong Kong; (2) CMIC, Shanghai Jiao Tong University, China; (3) Shanghai Innovation Institute, China

**Corresponding authors:** Yao, Jiangchao(Sunarker@sjtu.edu.cn); Han, Bo(bhanml@comp.hkbu.edu.hk)

**Source title:** arXiv

**Abbreviated source title:** arXiv

**Issue date:** October 3, 2025

**Publication year:** 2025

**Language:** English

**E-ISSN:** 23318422

**Document type:** Preprint (PP)

**Repository:** arXiv

**Abstract:** While reinforcement learning with verifiable rewards (RLVR) is effective to improve the reasoning ability of large language models (LLMs), its reliance on human-annotated labels leads to the scaling up dilemma, especially for complex tasks. Recent self-rewarding methods investigate a label-free alternative to unlock the reasoning capabilities of LLMs, yet they frequently encounter the non-negligible training collapse issue, as the single-view supervision signal easily forms the self-consistent illusion, yielding the reward hacking. Inspired by the success of self-supervised learning, we propose Co-rewarding, a novel self-supervised RL framework that improves training stability by seeking complementary supervision from another views. Specifically, we instantiate Co-rewarding in two ways: (1) Co-rewarding-I is a data-side instantiation that derives reward signals from contrastive agreement across semantically analogous questions; and (2) Co-rewarding-II is a model-side instantiation that maintains a slowly-updated reference teacher with pseudo labels to realize self-distillation. Intuitively, such instantiations introduce different levels of discrepancy to increase the difficulty of training collapse on trivial reasoning solutions. Empirically, Co-rewarding exhibits stable training across various setups, and outperforms other self-rewarding baselines by +3.31% improvements on average on multiple mathematical reasoning benchmarks, especially by +7.49% on Llama-3.2-3B-Instruct. Notably, Co-rewarding reaches or even surpasses RLVR with ground-truth (GT) label in several cases, such as a Pass@1 of 94.01% on GSM8K with Qwen3-8B-Base remarkably higher than GT. Our code is publicly available at <https://github.com/tmlr-group/Co-rewarding>. © 2025, CC BY.

**Number of references:** 0

**Main heading:** Personnel training

**Controlled terms:** Computational linguistics - Reinforcement learning - Self-supervised learning - Supervised learning - Teaching

**Uncontrolled terms:** Complex task - Ground truth - Label free - Language model - Reasoning ability - Reasoning capabilities - Reinforcement learnings - RL framework - Scaling-up - Two ways

**Classification code:** 901.2 Education - 912.3 Personnel - 1101.2 Machine Learning - 1101.2.1 Deep Learning - 1102.1 Computer Theory, Includes Computational Logic, Automata Theory, Switching Theory, Programming Theory

**Numerical data indexing:** Percentage 3.31E+00%, Percentage 7.49E+00%, Percentage 9.401E+01%

**DOI:** [10.48550/arXiv.2508.00410](https://doi.org/10.48550/arXiv.2508.00410)

**Compendex references:** YES

**Preprint ID:** 2508.00410v2

**Preprint source website:** <https://arxiv.org>

**Preprint ID type:** ARXIV

**Database:** Compendex

**Data Provider:** Engineering Village

Compilation and indexing terms, Copyright 2026 Elsevier Inc.